

Spacy and Custom NER

Josh Smith <kognate@gmail.com>

May 23, 2018

Outline

- 1 Natural Language Processing
- 2 Tools for NLP in Python
- 3 Demo of Basic Spacy
- 4 Named Entity Recognition in Spacy
- 5 Deploy the service

Stop Words

- These are glue words and things that can be ignored sometimes.

```
doc = nlp(('These are glue words and'  
          ' things that can be '  
          'ignored sometimes.'))  
[(x.lemma_, x.pos_, x.text)  
 for x  
 in [y  
     for y  
     in doc  
     if y.is_stop]]
```

be	VERB	are
and	CCONJ	and
that	ADJ	that
can	VERB	can
be	VERB	be
sometimes	ADV	sometimes

Parsing Human Languages

But you need a language model (at minimum).

Entity Recognition

This is knowing that Jason is a Person and not a Verb.

NLTK

This is one of the most mature and widely used frameworks for NLP

Spacy

This is the fastest. And has, IMHO, a great syntax

This is what we're going to talk about

Installs from pip

```
pip install spacy
```


don't forget to download at least one model

```
python -m spacy download en
```

First, we must import and load

```
import spacy
nlp = spacy.load('en')
```

Now we can use it

```
import spacy
nlp = spacy.load('en')
doc = nlp("Hello everyone, I've some good news to give you")
cleaned = [y for y
            in doc
            if not y.is_stop and y.pos_ != 'PUNCT']
raw = [(x.lemma_, x.pos_) for x in cleaned]
print(raw)
raw
```

hello	INTJ
-PRON-	PRON
have	VERB
good	ADJ
news	NOUN

```
import spacy
nlp = spacy.load('en')
to_analyze = ('Hello Code & Supply, '
              'my name is Josh and tonight '
              'we\'re in Pittsburgh')
doc = nlp(to_analyze)
ents = [(x.text, x.label_)
        for x in doc.ents]
print(ents)
ents
```

Josh	PERSON
tonight	TIME
Pittsburgh	GPE

gui tools:

- <https://prodi.gy/> from the creators of spacy

Train a model in Spacy

you can also create a vm for this:

```
gcloud compute instances create cans  
  --image-family ubuntu-1804-lts  
  --image-project ubuntu-os-cloud  
  --machine-type n1-highcpu-16
```

Some imports

```
import spacy
from spacy.matcher import PhraseMatcher
import plac
from pathlib import Path
import random
```

Utility function

This function converts the output of the **PhraseMatcher** to something usable in training. The training data needs a string index of characters (start, end, label) while the matched output uses index of words from an nlp document.

```
def offseter(lbl, doc, matchitem):  
    o_one = len(str(doc[0:matchitem[1]]))  
    subdoc = doc[matchitem[1]:matchitem[2]]  
    o_two = o_one + len(str(subdoc))  
    return (o_one, o_two, lbl)
```


Load and setup

Here we load `spacy` and setup the pipes for training.

```
nlp = spacy.load('en')

if 'ner' not in nlp.pipe_names:
    ner = nlp.create_pipe('ner')
    nlp.add_pipe(ner)
else:
    ner = nlp.get_pipe('ner')
```

Setup the phrase matches

This is to make our lives easier. Instead of setting this up by hand, we can use `PhraseMatcher` class from `spacy` to locate the text we want to label.

```
label = 'CIADIR'  
matcher = PhraseMatcher(nlp.vocab)  
for i in ['Gina Haspel', 'Gina', 'Haspel',]:  
    matcher.add(label, None, nlp(i))
```

What's that look like?

```
one = nlp('Gina Haspel was nominated in 2018')  
matches = matcher(one)  
[match for match in matches]
```

```
1.7539557946531887(+19) 0 1  
1.7539557946531887(+19) 0 2  
1.7539557946531887(+19) 1 2
```

Gather training data

```
res = []
to_train_ents = []
with open('gina_haspel.txt') as gh:
    line = True
    while line:
        line = gh.readline()
        mnlp_line = nlp(line)
        matches = matcher(mnlp_line)
        res = [offseter(label, mnlp_line, x)
                for x
                in matches]
        to_train_ents.append((line,
                              dict(entities=res)))
```

Actually Train The Recognizer

```
optimizer = nlp.begin_training()

other_pipes = [pipe
                for pipe
                in nlp.pipe_names
                if pipe != 'ner']

with nlp.disable_pipes(*other_pipes):  # only train NER
    for itn in range(20):
        losses = {}
        random.shuffle(to_train_ents)
        for item in to_train_ents:
            nlp.update([item[0]],
                      [item[1]],
                      sgd=optimizer,
                      drop=0.35,
                      losses=losses)
```

Build the Image

```
docker build -f Dockerfile-stemmer . -t gcr.io/codeandsupply
```

Gcloud

```
gcloud init
gcloud auth configure-docker
gcloud auth print-access-token
# then cut and paste that token
docker login -u oauth2accesstoken https://gcr.io
# you must have enabled the container registry before you can
docker push gcr.io/codeandsupply/stemmer:latest
gcloud container clusters create experimental-aone
gcloud container clusters get-credentials experimental-aone
kubectl run stemmer-server
    --image gcr.io/codeandsupply/stemmer:latest
    --port 8000
kubectl expose deployment stemmer-server --type "LoadBalancer"
kubectl get service stemmer-server
curl http://<HOSTNAME>:8000/

curl -s "http://127.0.0.1:8000/stemmer?source=Bill+is+a+nice"
    http://127.0.0.1:8000/stemmer?source=Bill+is+a+nice
```